



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

**Course - Distributed Computing (DC)**

<b>UID</b>	2023800080 2023800071 2023800079
<b>Name</b>	Jeet Patel Soham Padalia Arya Patel
<b>Class and Batch</b>	TE Computer Science and Engineering - Batch B1
<b>Date</b>	31/10/25
<b>Lab #</b>	9
<b>Aim</b>	To implement <b>Bully</b> and <b>Ring Election Algorithms</b> in a distributed system to elect a new coordinator when the current one fails.
<b>Objective</b>	<ol style="list-style-type: none"><li><i>To understand the concept of election algorithms in distributed systems.</i></li><li><i>To study how Bully and Ring algorithms select a new coordinator.</i></li><li><i>To simulate process coordination and failure recovery in a distributed environment.</i></li><li><i>To analyze message passing and process priority in leader election.</i></li><li><i>To compare the speed, fairness, and efficiency of both algorithms.</i></li></ol>
<b>Theory</b>	<p><i>In distributed systems, multiple processes work together without a central controller. When the process acting as a <b>coordinator (leader)</b> fails, an <b>election algorithm</b> is used to select a new one.</i></p> <p><b>1. Bully Algorithm:</b></p> <ul style="list-style-type: none"><li><i>The process with the highest ID becomes the coordinator.</i></li><li><i>When a process detects the leader's failure, it sends an <b>ELECTION</b> message to all higher-ID processes.</i></li><li><i>If no higher process replies, it declares itself as the coordinator and informs others.</i></li></ul> <p><b>2. Ring Algorithm:</b></p> <ul style="list-style-type: none"><li><i>All processes are arranged in a <b>logical ring</b>.</i></li><li><i>When a coordinator fails, the process noticing it sends an <b>ELECTION</b> message around the ring.</i></li><li><i>Each process adds its ID, and the one with the highest ID is chosen as the new coordinator once the message completes the ring.</i></li></ul> <p><i>These algorithms ensure <b>fault tolerance, coordination, and continuity</b> in distributed systems.</i></p>
<b>Procedure:</b>	<p><b>Bully:-</b></p> <ol style="list-style-type: none"><li><i>Assign a unique ID to each process (higher ID = higher priority).</i></li><li><i>Start all processes and elect an initial coordinator (highest ID).</i></li></ol>



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

- |  |  |
|--|--|
|  | <ol style="list-style-type: none"><li>3. When a process detects the coordinator has failed, it sends an <i>ELECTION</i> message to all processes with higher IDs.</li><li>4. If no higher process responds, the initiating process becomes the new coordinator and broadcasts a <i>COORDINATOR</i> message to all nodes.</li><li>5. If any higher process replies with <i>OK</i>, that higher process takes over the election.</li><li>6. Observe the output to confirm that the highest active process becomes the new coordinator.</li></ol> |
|--|--|

**Procedure for Ring Election Algorithm:**

- |  |  |
|--|--|
|  | <ol style="list-style-type: none"><li>1. Arrange all processes in a logical ring and assign each a unique ID.</li><li>2. When a process detects that the coordinator has failed, it initiates an <i>ELECTION</i> message containing its ID and sends it to the next process in the ring.</li><li>3. Each process adds its own ID to the message and forwards it to the next node.</li><li>4. When the message completes the ring, the process with the highest ID is chosen as the new coordinator.</li><li>5. The coordinator sends a <i>COORDINATOR</i> message around the ring to inform all nodes.</li><li>6. Verify that all nodes recognize the same coordinator after election.</li></ol> |
|--|--|

**Implementation / Code 1**

	<pre>1A: bully.json: [     { "pid": 1, "host": "127.0.0.1", "port": 6001 },     { "pid": 2, "host": "127.0.0.1", "port": 6002 },     { "pid": 3, "host": "127.0.0.1", "port": 6003 } ] bully_node.py import argparse import json import socket import threading import time  # Load nodes from json with open('nodes.json') as f:     nodes = json.load(f)  def get_node(pid):     return next(n for n in nodes if n["pid"] == pid)  class BullyNode:     def __init__(self, pid):         self.pid = pid         self.node_info = get_node(pid)         self.coordinator = None         self.alive = True         self.lock = threading.Lock()</pre>
--	---



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

```
def start(self):
    # Start server thread
    threading.Thread(target=self.server_thread, daemon=True).start()
    time.sleep(2) # Wait for all nodes to start
    self.start_election()

    # Periodically check coordinator alive
    while self.alive:
        time.sleep(5)
        if self.coordinator is None or not self.ping(self.coordinator):
            print(f"[{self.pid}] Coordinator {self.coordinator} not responding -> starting election")
            self.start_election()

def server_thread(self):
    s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
    s.bind((self.node_info['host'], self.node_info['port']))
    s.listen(5)
    print(f"[{self.pid}] Listening on {self.node_info['host']}:{self.node_info['port']}")

    while self.alive:
        conn, addr = s.accept()
        threading.Thread(target=self.handle_connection, args=(conn,), daemon=True).start()

def handle_connection(self, conn):
    msg = conn.recv(1024).decode()
    if msg.startswith("ELECTION"):
        sender_pid = int(msg.split()[1])
        print(f"[{self.pid}] Received ELECTION from {sender_pid}")
        # Reply OK if self.pid > sender_pid
        if self.pid > sender_pid:
            conn.sendall("OK".encode())
            # Start election
            self.start_election()
        else:
            # Ignore
            pass
    elif msg.startswith("COORDINATOR"):
        new_coord = int(msg.split()[1])
        with self.lock:
            self.coordinator = new_coord
        print(f"[{self.pid}] Received COORDINATOR message: new coordinator = {new_coord}")
        conn.close()
```



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

```
def start_election(self):
    print(f"[{self.pid}] Initiating election...")
    higher_nodes = [n for n in nodes if n['pid'] > self.pid]
    received_ok = False
    for node in higher_nodes:
        if self.send_message(node, f"ELECTION {self.pid}"):
            print(f"[{self.pid}] Received OK from {node['pid']}")"
            received_ok = True

    if not received_ok:
        # I am the coordinator
        with self.lock:
            self.coordinator = self.pid
            print(f"[{self.pid}] I am the new COORDINATOR")
            # Broadcast coordinator message
            for node in nodes:
                if node['pid'] != self.pid:
                    self.send_message(node, f"COORDINATOR {self.pid}")

def send_message(self, node, message):
    try:
        s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
        s.settimeout(2)
        s.connect((node['host'], node['port']))
        s.sendall(message.encode())
        if message.startswith("ELECTION"):
            response = s.recv(1024).decode()
            if response == "OK":
                s.close()
                return True
            s.close()
    except:
        return False
    return False

def ping(self, pid):
    node = get_node(pid)
    try:
        s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
        s.settimeout(2)
        s.connect((node['host'], node['port']))
        s.close()
        return True
    except:
        return False
```



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

	<pre>if __name__ == "__main__":     parser = argparse.ArgumentParser()     parser.add_argument("--pid", type=int, required=True)     args = parser.parse_args()      node = BullyNode(args.pid)     node.start()</pre>
<b>Output 1</b>	<pre>students@cse-406-b-OptiPlex-SFF-7010:~/Desktop/kt\$ python3 bully_node.py --pid 1 [1] Listening on 127.0.0.1:6001 [1] Initiating election... [1] I am the new COORDINATOR [1] Received COORDINATOR message: new coordinator = 2 [1] Received COORDINATOR message: new coordinator = 3  students@cse-406-b-OptiPlex-SFF-7010:~/Desktop/kt\$ python3 bully_node.py --pid 2 [2] Listening on 127.0.0.1:6002 [2] Initiating election... [2] I am the new COORDINATOR [2] Received COORDINATOR message: new coordinator = 3  students@cse-406-b-OptiPlex-SFF-7010:~/Desktop/kt\$ python3 bully_node.py --pid 3 [3] Listening on 127.0.0.1:6003 [3] Initiating election... [3] I am the new COORDINATOR</pre>
<b>Implementation/ Code 2</b>	<pre>1B: import argparse import json import socket import threading import time  with open('nodes.json') as f:     nodes = json.load(f)  def get_node(pid):</pre>



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

```
return next(n for n in nodes if n['pid'] == pid)

class RingNode:
    def __init__(self, pid):
        self.pid = pid
        self.node_info = get_node(pid)
        self.next_node = self.get_next_node()
        self.coordinator = None
        self.alive = True
        self.lock = threading.Lock()

    def get_next_node(self):
        idx = next(i for i, n in enumerate(nodes) if n['pid'] == self.pid)
        return nodes[(idx + 1) % len(nodes)]

    def start(self):
        threading.Thread(target=self.server_thread, daemon=True).start()
        time.sleep(5) # Increased sleep to give time for all nodes to start

        self.start_election()

        while self.alive:
            time.sleep(5)
            if self.coordinator is None or not self.ping(self.coordinator):
                print(f"[{self.pid}] Coordinator {self.coordinator} not responding -> starting election")
                self.start_election()

    def server_thread(self):
        s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
        s.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
        s.bind((self.node_info['host'], self.node_info['port']))
        s.listen(5)
        print(f"[{self.pid}] Listening on {self.node_info['host']}:{self.node_info['port']} | next -> {self.next_node['pid']} @ {self.next_node['host']}:{self.next_node['port']}")

        while self.alive:
            try:
                conn, addr = s.accept()
                threading.Thread(target=self.handle_connection, args=(conn,), daemon=True).start()
            except Exception as e:
                print(f"[{self.pid}] Server accept error: {e}")

    def handle_connection(self, conn):
        try:
```



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

```
msg = conn.recv(1024).decode()
print(f"[{self.pid}] Received message: {msg}")
if msg.startswith("ELECTION"):
    data = msg.split(maxsplit=2)
    origin = int(data[1])
    if len(data) > 2:
        ids_str = data[2].strip("[]")
        if ids_str:
            ids = list(map(int, ids_str.split(",")))
        else:
            ids = []
    else:
        ids = []

    if self.pid not in ids:
        ids.append(self.pid)

    if origin == self.pid:
        # Election message returned to origin, election done
        coordinator = max(ids)
        print(f"[{self.pid}] ELECTION result -> coordinator = {coordinator}")
        with self.lock:
            self.coordinator = coordinator
            self.send_message(self.next_node, f"COORDINATOR {coordinator}")
    else:
        print(f"[{self.pid}] Forwarding ELECTION from origin={origin} ids={ids}")
        self.send_message(self.next_node, f"ELECTION {origin} [{','.join(map(str, ids))}]")

elif msg.startswith("COORDINATOR"):
    new_coord = int(msg.split()[1])
    with self.lock:
        self.coordinator = new_coord
    print(f"[{self.pid}] Received COORDINATOR announcement: {new_coord}")
    if new_coord != self.pid:
        self.send_message(self.next_node, msg)
except Exception as e:
    print(f"[{self.pid}] Error handling connection: {e}")
finally:
    conn.close()

def start_election(self):
    print(f"[{self.pid}] Initiating election...")
    sent = self.send_message(self.next_node, f"ELECTION {self.pid} [{self.pid}]")
    if not sent:
        print(f"[{self.pid}] Failed to send election message to next node")
```



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

```
{self.next_node['pid']}")\n\n    def send_message(self, node, message):\n        try:\n            s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)\n            s.settimeout(2)\n            s.connect((node['host'], node['port']))\n            s.sendall(message.encode())\n            s.close()\n            print(f"[{self.pid}] Sent message to {node['pid']}: {message}")\n            return True\n        except Exception as e:\n            print(f"[{self.pid}] Failed to send message to {node['pid']}\n({node['host']},{node['port']}): {e}")\n            return False\n\n    def ping(self, pid):\n        node = get_node(pid)\n        try:\n            s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)\n            s.settimeout(2)\n            s.connect((node['host'], node['port']))\n            s.close()\n            return True\n        except Exception as e:\n            print(f"[{self.pid}] Ping failed to {pid}: {e}")\n            return False\n\nif __name__ == "__main__":\n    parser = argparse.ArgumentParser()\n    parser.add_argument("--pid", type=int, required=True)\n    args = parser.parse_args()\n\n    node = RingNode(args.pid)\n    node.start()
```



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

**Output 2**

```
students@cse-406-b-OptiPlex-SFF-7010:~/Desktop/kt$ python3 ring.py --pid 1
[1] Listening on 127.0.0.1:6001 | next -> 2 @ 127.0.0.1:6002
[1] Initiating election...
[1] Sent message to 2: ELECTION 1 [1]
[1] Coordinator None not responding -> starting election
[1] Initiating election...
[1] Sent message to 2: ELECTION 1 [1]
[1] Received message: ELECTION 2 [2,3]
[1] Forwarding ELECTION from origin=2 ids=[2, 3, 1]
[1] Sent message to 2: ELECTION 2 [2,3,1]
[1] Coordinator None not responding -> starting election
[1] Initiating election...
[1] Sent message to 2: ELECTION 1 [1]
[1] Received message: ELECTION 1 [1,2,3]
[1] ELECTION result -> coordinator = 3
[1] Sent message to 2: COORDINATOR 3
[1] Received message: ELECTION 3 [3]
[1] Forwarding ELECTION from origin=3 ids=[3, 1]
[1] Sent message to 2: ELECTION 3 [3,1]
[1] Received message: COORDINATOR 3
[1] Received COORDINATOR announcement: 3
[1] Sent message to 2: COORDINATOR 3

students@cse-406-b-OptiPlex-SFF-7010:~/Desktop/kt$ python3 ring.py --pid 2
[2] Listening on 127.0.0.1:6002 | next -> 3 @ 127.0.0.1:6003
[2] Received message: ELECTION 1 [1]
[2] Forwarding ELECTION from origin=1 ids=[1, 2]
[2] Failed to send message to 3 (127.0.0.1:6003): [Errno 111] Connection refused
[2] Initiating election...
[2] Failed to send message to 3 (127.0.0.1:6003): [Errno 111] Connection refused
[2] Failed to send election message to next node 3
[2] Received message: ELECTION 1 [1]
[2] Forwarding ELECTION from origin=1 ids=[1, 2]
[2] Failed to send message to 3 (127.0.0.1:6003): [Errno 111] Connection refused
[2] Coordinator None not responding -> starting election
[2] Initiating election...
[2] Sent message to 3: ELECTION 2 [2]
[2] Received message: ELECTION 2 [2,3,1]
[2] ELECTION result -> coordinator = 3
[2] Sent message to 3: COORDINATOR 3
[2] Received message: ELECTION 1 [1]
[2] Forwarding ELECTION from origin=1 ids=[1, 2]
[2] Sent message to 3: ELECTION 1 [1,2]
[2] Received message: COORDINATOR 3
[2] Received COORDINATOR announcement: 3
[2] Sent message to 3: COORDINATOR 3
[2] Received message: ELECTION 3 [3,1]
[2] Forwarding ELECTION from origin=3 ids=[3, 1, 2]
[2] Sent message to 3: ELECTION 3 [3,1,2]
[2] Received message: COORDINATOR 3
[2] Received COORDINATOR announcement: 3
[2] Sent message to 3: COORDINATOR 3
```



**BHARATIYA VIDYA BHAVAN'S**  
**SARDAR PATEL INSTITUTE OF TECHNOLOGY**  
(Empowered Autonomous Institute Affiliated to University of Mumbai)  
[Knowledge is Nectar]

**Department of Computer Engineering**

```
students@cse-406-b-OptiPlex-SFF-7010:~/Desktop/k$ python3 ring.py --pid 3
[3] Listening on 127.0.0.1:6003 | next -> 1 @ 127.0.0.1:6001
[3] Received message: ELECTION 2 [2]
[3] Forwarding ELECTION from origin=2 ids=[2, 3]
[3] Sent message to 1: ELECTION 2 [2,3]
[3] Received message: COORDINATOR 3
[3] Received COORDINATOR announcement: 3
[3] Received message: ELECTION 1 [1,2]
[3] Forwarding ELECTION from origin=1 ids=[1, 2, 3]
[3] Sent message to 1: ELECTION 1 [1,2,3]
[3] Received message: COORDINATOR 3
[3] Received COORDINATOR announcement: 3
[3] Initiating election...
[3] Sent message to 1: ELECTION 3 [3]
[3] Received message: ELECTION 3 [3,1,2]
[3] ELECTION result -> coordinator = 3
[3] Sent message to 1: COORDINATOR 3
[3] Received message: COORDINATOR 3
[3] Received COORDINATOR announcement: 3
[3] Received message:
[3] Received message:
[3] Received message:
```

<b>Conclusion</b>	<p>From this experiment, I learned how Bully and Ring Election Algorithms work in distributed systems to select a new coordinator when the current one fails. Both algorithms ensure that the system continues to function even after a node crash, which is important for reliability and fault tolerance.</p> <p>In the Bully Algorithm, the process with the highest ID quickly takes over as the coordinator, making it faster but with higher message overhead.</p> <p>In contrast, the Ring Algorithm provides a more balanced and fair approach since every process gets a chance to participate, though it can take more time as messages circulate through the ring.</p> <p>Overall, this lab helped me understand how leader election, process coordination, and failure recovery are managed practically in distributed systems. It also improved my understanding of inter-process communication and how algorithms maintain system stability after a failure.</p>
<b>References</b>	<ol style="list-style-type: none"><li>1. A. S. Tanenbaum and M. Van Steen, <i>Distributed Systems: Principles and Paradigms</i>, 2nd Edition, Pearson Education, 2007.</li><li>2. G. Coulouris, J. Dollimore, T. Kindberg, <i>Distributed Systems: Concepts and Design</i>, 5th Edition, Pearson, 2012.</li><li>3. Course Lab Manual – DC Experiment 9: <i>Bully and Ring Election Algorithms</i>.</li><li>4. Online reference: GeeksforGeeks – <i>Bully and Ring Election Algorithms in Distributed Systems</i> (<a href="https://www.geeksforgeeks.org/">https://www.geeksforgeeks.org/</a>).</li></ol>